

Deepak Mallampati

MACHINE LEARNING ENGINEER - HEALTHCARE AI

(330) 554-0770 | deepakmallampati00@gmail.com | linkedin.com/in/deepak-mallampati |
US-based, F-1 OPT, open to relocation

PROFESSIONAL SUMMARY

ML Engineer with 4+ years building AI systems for regulated healthcare. Deep experience with privacy-preserving clinical ML (re-identification risk cut from 3.38% to 0.32% with 99% predictive performance retained), HIPAA-compliant data workflows, and multimodal LLM fine-tuning (QLoRA on H100). Works in PyTorch and GCP Vertex AI across healthcare NLP, RAG, and video/image pipelines. Comfortable owning model training, evaluation, and clinical-software integration.

CORE COMPETENCIES & SKILLS

PyTorch • Multimodal LLM Fine-tuning • QLoRA / LoRA / PEFT • Healthcare AI & HIPAA • Privacy-Preserving ML • Medical Data Pipelines • GCP Vertex AI • Model Optimization • Clinical Integration
Generative AI & LLMs: GPT-4o, Claude Sonnet, Gemini, LLaMA | LangChain/LangGraph | RAG Architecture | Multi-Agent Systems | Fine-tuning (LoRA, QLoRA, PEFT) | Hugging Face Transformers
ML & NLP: PyTorch, TensorFlow, Keras | NER, Text Classification, Summarization | Model Optimization | RAGAS Evaluation | BERTScore | Drift Detection
Vector Search & Retrieval: FAISS, Pinecone, Weaviate | Semantic & Hybrid Search | Cross-Encoder Re-ranking | Semantic Caching
MLOps & Cloud: AWS (Bedrock, SageMaker), Azure OpenAI, GCP Vertex AI | Docker, Kubernetes | Terraform | CI/CD | MLflow, Weights & Biases
Data & Compliance: PostgreSQL, MongoDB, Redis | Spark, Kafka | HIPAA, SOC 2, PCI DSS | PII Masking | Prompt Injection Mitigation
Languages & Tools: Python, JavaScript/TypeScript, Java, SQL | FastAPI, React | Git, Jira

PROFESSIONAL EXPERIENCE

Progress Solutions

Jul 2025 - Present

AI Engineer

USA

- Architected production-grade **agentic LLM systems** on a conductor-subagent orchestration framework, decomposing complex objectives into context-scoped subagents that execute autonomously - reducing token consumption by **42%** while sustaining task accuracy across multi-step workflows.
- Engineered high-performance **retrieval-augmented generation (RAG)** pipelines over large-scale healthcare document corpora, combining dense vector retrieval with cross-encoder re-ranking to lift answer relevance and materially reduce hallucination on domain-specific queries.
- Built a structured **LLM evaluation and monitoring framework** (RAGAS, BERTScore, custom domain metrics) with latency and accuracy dashboards to benchmark model iterations pre-deployment and surface regressions before release.
- Optimized inference cost and latency through prompt compression, semantic caching, and model routing, sustaining SLA targets while reducing per-query overhead.

Vanguard

Jan 2024 - Jan 2025

AI Engineer Intern

USA

- Developed **AI agents and backend services** for Vanguard's internal AI platform across **25+ AI agents** and workflows, integrating with production data pipelines and observability tooling.
- Optimized LLM prompts and evaluated **GPT-4o, Claude Sonnet, and Gemini**, improving task success rates by **27%** and informing platform model-selection decisions.
- Built **12 APIs and microservices** powering new AI capabilities, cutting p95 latency from 1.5s to **800ms** and supporting **100,000+ requests daily** across 20+ internal teams.
- Engineered content controls for secure, policy-compliant AI responses, reducing unsafe outputs by **42%** while maintaining response quality.

Kent State University

Oct 2023 - Jan 2024

ML & Gen AI Research Assistant

USA

- Built a **privacy-preserving ML pipeline** on a clinical hospital-readmission dataset, applying k-anonymity, l-diversity, and Laplace differential privacy across 8 clinical features.
- Reduced record-linkage re-identification risk from **3.38% to 0.32%** while retaining **99% of baseline predictive performance** (Random Forest AUC 0.689 vs. 0.696).
- Fine-tuned **Qwen3-27B via 4-bit QLoRA** on an NVIDIA H100 with a curated instruction dataset for controllable long-form generation.

Emerson

Jun 2022 - Apr 2023

ML Trainee

India

- Trained supervised **regression and classification models** on operational pipeline-storage data for equipment-failure prediction, anomaly detection, and capacity forecasting.
- Fine-tuned **BERT and RoBERTa** for entity extraction and text classification on domain-specific datasets, achieving **89% F1-score** on custom NER tasks.

PROJECTS

MangaLens [Chrome Web Store](#)

Shipped Chrome extension delivering real-time AI translation of manga and webtoons across 29+ sites, turning a multi-step manual translation chore into a single inline action.

Privacy-Preserving Clinical ML Pipeline

Framework combining k-anonymity, l-diversity, and differential privacy so researchers can analyze sensitive patient data without exposing identities, quantifying the privacy-utility trade-off precisely.

Story Director LLM

Fine-tuned language model that generates structured short-form video narratives from a single prompt, paired with an automated rendering pipeline.