Deepak Mallampati

+1-330-554-0770 | deepakmallampati00@gmail.com | linkedin.com/in/deepak-mallampati | github.com/Deepak0070 |
Kent, OH - US-based, F-1 OPT, open to sponsorship

PROFESSIONAL SUMMARY

Applied AI engineer (3.5+ years) with RAG as primary production specialty. Healthcare RAG with recursive semantic chunking + transformer embeddings: ~35% retrieval precision gain, >30% irrelevant retrieval reduction, >90% retrieval-grounded alignment in evaluation. Suvidha document Q&A with RAG: hallucinations -30%, contextual accuracy >85%. Multi-step agentic workflows + schema-validated JSON contracts (Agentic Healthcare Claims). LangChain + LlamaIndex in production stack. Solo shipped Manga Lens to Chrome Web Store. Ready to embed with Tavily customers to build RAG pipelines on top of the search engine for AI agents.

CORE COMPETENCIES

Applied AI / LLMs • RAG + Grounding • Agentic Workflows • Schema Contracts • Eval Harnesses • Python / FastAPI • TypeScript / React • Docker / CI-CD • Multimodal Vision • Predictive ML

PROFESSIONAL EXPERIENCE

Progress Solutions Inc. *Jul 2025 - Present*

Jr. AI/ML Engineer Trainee - Agent Engineering, Production LLM

USA - Healthcare Technology, AI Consulting

- Built **Retrieval-Augmented Generation (RAG)** for clinical knowledge retrieval; recursive semantic chunking and transformer embeddings improved retrieval precision **~35%**, reduced irrelevant retrieval **>30%**, response alignment **>90%**.
- Developed **agentic LLM workflows** for multi-step healthcare queries with tool discipline, structured reasoning, and grounding rules; response stability **~25%**, hallucinations cut **>30%**.
- Shipped **predictive ML pipelines** (scikit-learn, XGBoost) for patient no-show and care engagement scoring; recall **+15-20%** on high-risk categories, precision held **>90%** via class weighting, stratified sampling, threshold calibration.
- Packaged ML/LLM inference as **FastAPI / Flask** services on **Docker** with structured logging and load simulation; post-deploy defects **~30%**.
- EHR preprocessing (Pandas, NumPy); dataset reliability **>98%**, downstream instability **-40%**.
- HIPAA-conscious governance: de-identification, data lineage, evaluation audit trails, stakeholder-facing system-limitation docs.

Energy Solutions International (Emerson) *Jun 2022 - Apr 2023*

Database & DevOps Performance Engineer (Intern)

Hyderabad, India - Oil & Gas ERP

- Optimized **T-SQL stored procedures** on a compliance-sensitive ERP (contracts, nominations, allocations, invoicing); query time **-20%**, retrieval latency **-25%**.
- Built **Jenkins CI/CD pipelines** for schema updates and stored-procedure deployments; deploy errors **-30%**, release cycle **-35 to -40%**.
- Designed performance dashboards with **SQL Server DMVs + Grafana**; incident recurrence **-25%**, root-cause resolution **-30%**.
- Tuned index maintenance and partition-aware queries; reconciliation runtime **-15%**, deadlocks **-15%**.
- Supported RBAC and audit logging for financial modules in a regulated environment.

Suvidha Foundation *Jan 2022 - Mar 2022*

Machine Learning Engineer

Hyderabad, India - Nonprofit, Video Analytics, Applied NLP

- Built **transformer-based video summarization** over 5,000+ recorded sessions; hierarchical summarization and timestamp-aligned clip extraction cut manual review **60-70%**.
- AI-selected highlights aligned with human-curated moments **~85%**.
- Implemented **document Q&A with RAG**; hallucinations **~30%**, contextual accuracy **>85%**.
- Deployed as **Flask** API with a lightweight web interface for non-technical staff.

PROJECTS

Agentic Healthcare Claims Processing & Fraud Risk Intelligence System

5-stage multi-agent pipeline (Intake, Validation, Consistency, Duplicate, Risk Scoring) with schema-validated JSON contracts between agents to prevent cascading hallucinations. RAG-grounded CPT/ICD validation against vector-indexed policies. Audit-ready explainable risk scoring.

Python, LangChain, FastAPI, pgvector, GPT-4, schema contracts

Manga Lens - Multi-Provider AI Vision Chrome Extension

Real-time manga translation extension shipped to the Chrome Web Store. Multi-provider abstraction (Claude Sonnet, GPT-4o mini, GPT-4.1 Nano, Ollama) with per-provider payload handling for worst-case failure isolation. 29 sites supported across Manifest V3 service workers.

TypeScript, Manifest V3, Service Workers, Claude / OpenAI / Ollama

Dream Decoder - Multimodal Creative Platform

End-to-end multimodal pipeline: text description, intermediate prompt transformation, grounded image generation. ~30% contextual alignment gain, ~25-30% first-pass image success.

FastAPI, React + Vite, Tailwind, Perplexity Sonar, Replicate

YOLOv8 Driver Drowsiness Detection

Real-time computer vision for driver drowsiness; sub-50ms inference on edge hardware. Walk-forward validation and class-weighted training for imbalanced safety data.

PyTorch, YOLOv8, OpenCV, ONNX

EDUCATION

Master of Science, Computer Science - *Kent State University* *Aug 2023 - May 2025*

Focus: Applied Machine Learning, NLP, Database Systems.

Bachelor of Technology, Computer Science - *KL University* *Jun 2019 - May 2023*

Founder / Lead of E-Farming platform (university project, pilot with 3 cooperatives).

CERTIFICATIONS

Deep Learning Specialization - *DeepLearning.AI (Coursera)* *2024*

SKILLS

Languages: Python, TypeScript, JavaScript, SQL (T-SQL, PostgreSQL, SQLite), Java, C++

AI / ML: PyTorch, scikit-learn, XGBoost, LangChain, LlamaIndex, OpenAI, Anthropic, Hugging Face Transformers, Diffusers, YOLOv8

LLM / GenAI: RAG, agentic workflows, structured outputs, prompt engineering, evaluation pipelines, guardrails, grounding, embeddings, vector search

Backend: FastAPI, Flask, REST, WebSockets, Docker, Celery, Playwright

Frontend: React, Next.js, Vite, Tailwind, Chrome Extensions (Manifest V3)

Data: PostgreSQL + pgvector, MongoDB, Redis, Pandas, NumPy, Snowflake

DevOps: Docker, Jenkins, GitHub Actions, Grafana, Prometheus, AWS (EC2, S3, Lambda)

Domains: Healthcare AI (HIPAA-conscious), Applied AI, Enterprise AI, Workflow Automation, Computer Vision, Multimodal