

Deepak Mallampati

(330) 554-0770 | deepakmallampati00@gmail.com | linkedin.com/in/deepak-mallampati | github.com/deepak-glitch |

Kent, OH | US-based, F-1 OPT (open to sponsorship)

PROFESSIONAL SUMMARY

Applied AI engineer (2.5+ years) shipping production LLM systems: clinical RAG (~35% precision lift, >30% hallucination reduction), multi-agent healthcare claims pipeline with JSON-contracted agents and audit-ready traces, FastAPI/Docker packaging for HIPAA-conscious deployment. Comfortable embedding with customer teams - fast prototyping, requirement translation, end-to-end delivery - and grounding LLM behaviors with retrieval, evaluation, and tool discipline.

CORE COMPETENCIES

Applied AI • Forward Deployed • Agentic Workflows • RAG • Customer-Facing Delivery • Prototyping • FastAPI • Docker • Python • LLM Apps • Production Reliability • Evals

PROFESSIONAL EXPERIENCE

Jr. AI/ML Engineer Trainee

Progress Solutions Inc. · USA · Jul 2025 - Present

Built **Retrieval-Augmented Generation (RAG)** systems for clinical knowledge retrieval; recursive semantic chunking + transformer embeddings improved retrieval precision ~**35%** and cut irrelevant retrieval >**30%**; retrieval-grounded response alignment exceeded **90%**.

Developed **agentic LLM workflows** (eligibility checks, care-navigation, documentation clarification) with structured reasoning, tool discipline, and grounding rules; improved response stability ~**25%**.

Shipped **predictive ML pipelines** (scikit-learn, XGBoost) for patient no-show prediction, care engagement scoring, and support prioritization; improved recall **15-20%** on high-risk categories while holding precision >**90%**.

Packaged ML/LLM inference as **FastAPI/Flask** REST services, containerized with **Docker**, structured logs + load simulation; reduced post-deployment defects ~**30%**.

Built preprocessing pipelines (Pandas, NumPy) for EHR extracts, appointment histories, support tickets; raised dataset reliability above **98%** and cut downstream model instability ~**40%**.

Maintained HIPAA-conscious data governance: de-identification, lineage documentation, evaluation audit trails, stakeholder-facing system-limitation docs.

Database & DevOps Performance Engineer (Intern)

Energy Solutions International (Emerson) · Hyderabad, India · Jun 2022 - Apr 2023

Optimized SQL and **T-SQL stored procedures** for the Synthesis Order-to-Cash platform; cut query execution time ~**20%** and data retrieval latency ~**25%**.

Built **CI/CD pipelines with Jenkins** for schema and stored-procedure deployments; reduced deployment errors >**30%** and release cycle time ~**35-40%**.

Designed performance dashboards (SQL Server DMVs, Grafana) for query/deadlock/CPU-I/O monitoring; cut incident recurrence ~**25%** and improved root-cause speed ~**30%**.

Tuned indexing, statistics, and reconciliation queries; improved reconciliation runtime ~**15%**; supported RBAC + audit logging for compliance-sensitive financial modules.

Machine Learning Engineer

Suvidha Foundation · Hyderabad, India · Jan 2022 - Mar 2022

Built a **transformer-based video summarization system** across 5,000+ recorded sessions; hierarchical summarization + timestamp-aligned clip extraction reduced manual review **60-70%**.

Achieved ~**85% alignment** with human-curated highlights; implemented document Q&A with RAG (semantic chunking + embeddings); cut hallucinated outputs ~**30%** and raised contextual accuracy >**85%**.

Deployed the stack as a **Flask-based API** with a lightweight web UI for non-technical staff.

SELECTED PROJECTS

Agentic Healthcare Claims Processing & Fraud Risk Intelligence

Multi-agent pipeline (Intake Normalization → Validation → Consistency Analysis → Duplicate Detection → Fraud Risk Scoring) with *schema-validated JSON contracts between agents* to prevent cascading hallucinations. RAG-grounded CPT/ICD validation against vector-indexed policy documents; approximate-nearest-neighbor similarity search for duplicate claim detection; explainable risk scoring with audit-ready reasoning traces.

Manga Lens — Chrome Extension for Real-Time AI Manga/Webtoon Translation (*shipped, Chrome Web Store*)

Full-stack browser extension built independently (Manifest V3, content scripts, service workers, captureVisibleTab). Integrates four AI vision providers (Claude Sonnet, GPT-4o mini, GPT-4.1 Nano, Ollama). Multi-section capture pipeline handles tall webtoon panels via viewport-slice screenshots with coordinate remapping and dedup. Seven-day translation cache, per-domain selector configs for 29 manga/webtoon sites, narrowed host permissions.

Dream Decoder — AI-Powered Multimodal Creative Intelligence Platform

Full-stack FastAPI + React/TypeScript/Vite/Tailwind app that interprets dreams, generates poetic reinterpretations, and synthesizes dreamscapes via coordinated multimodal APIs (Perplexity Sonar, Replicate image models). Introduced *intermediate structured prompt transformation layers* — improved contextual alignment ~30% and first-pass image success ~25-30% over naive direct-prompt orchestration.

Agentic Pixel Character Synthesis & Animation Engine (*ongoing*)

Phased generative AI system for identity-consistent, pose-controlled pixel character synthesis with sprite-sheet export. Stable Diffusion fine-tuning + LoRA for identity consistency, ControlNet + OpenPose/MediaPipe for pose conditioning, latent-space consistency for animation smoothness, and an *LLM-based agent orchestrator* that decomposes high-level prompts into generation tasks. PyTorch, Hugging Face Diffusers, FastAPI, Docker.

Driver Drowsiness Detection — Real-Time YOLOv8 Fatigue Monitoring

Replaced two-stage CNN pipeline with unified YOLOv8 detection-and-classification model; reduced inference latency ~30%. LabelImg annotations, augmentation for lighting/head-pose robustness, sliding-window confidence aggregation suppressed blink-driven false positives (~25% reduction), adaptive frame skipping, and NMS tuning for stable real-time operation via OpenCV.

Video Summarization & Highlight Generation System (Suvidha Foundation)

Transformer-based hierarchical abstractive summarization mapped back to source timestamps for automated clip extraction. *60-70% reduction in manual review time*, ~85% highlight alignment with human curation, Flask-based API for non-technical users.

IoT + ML Smart Building Indoor Temperature Prediction System

Arduino + DHT11 + DS3231 sensor pipeline with PLX-DAQ logging; time-series forecasting with walk-forward validation comparing Random Forest, SVR, and Linear Regression. Lag-window feature engineering ($n_{in}=6$); Random Forest achieved ~20-30% lower MAE vs linear baseline.

EDUCATION

Master's Degree — Kent State University

Ohio, USA · Aug 2023 - May 2025

CERTIFICATIONS

Open to specialization-specific certifications on request.

SKILLS

Languages & Frameworks: Python, FastAPI, Flask, SQL (T-SQL, PostgreSQL, SQLite), TypeScript, React, C++ (Arduino), PHP

LLM & GenAI: LLM apps, RAG, agentic workflows, structured outputs, prompt engineering, evaluation pipelines, guardrails, grounding, embeddings, vector search, responsible AI

ML & Multimodal: scikit-learn, XGBoost, PyTorch, Hugging Face Transformers/Diffusers, LangChain, LlamaIndex, YOLOv8, OpenCV, ControlNet, OpenPose/MediaPipe, Stable Diffusion, video analytics

Data & Infra: Pandas, NumPy, Docker, Jenkins, Grafana, CI/CD, REST APIs, cloud-ready deployment, observability/logging

Domains: Healthcare AI (HIPAA-conscious), Applied AI, Enterprise AI, Workflow Automation, Computer Vision, Multimodal Systems