

Deepak Mallampati

+1-330-554-0770 | deepakmallampati00@gmail.com | linkedin.com/in/deepak-mallampati | github.com/Deepak0070 |
Kent, OH

PROFESSIONAL SUMMARY

Applied AI engineer with diffusion + multimodal production work - Stable Diffusion + LoRA + ControlNet identity-consistent character synthesis pipeline (Diffusers, PyTorch, LLM agent orchestrator) and a Chrome Web Store extension integrating four AI vision providers under a single adapter layer. Also shipped a YOLOv8 detection-classification pipeline with ~30% inference-latency reduction, a clinical RAG system with ~35% retrieval-precision gain in a HIPAA-conscious environment, and FastAPI/Docker inference services with ~30% post-deploy defect reduction. US-based, F-1 OPT.

CORE COMPETENCIES

Applied AI / LLMs • RAG + Grounding • Agentic Workflows • Schema Contracts • Eval Harnesses • Python / FastAPI • TypeScript / React • Docker / CI-CD • Multimodal Vision • Predictive ML

PROFESSIONAL EXPERIENCE

Progress Solutions Inc.

Jul 2025 - Present

Jr. AI/ML Engineer Trainee - Agent Engineering, Production LLM

USA - Healthcare Technology, AI Consulting

- Built **Retrieval-Augmented Generation (RAG)** for clinical knowledge retrieval; recursive semantic chunking and transformer embeddings improved retrieval precision ~**35%**, reduced irrelevant retrieval >**30%**, response alignment >**90%**.
- Developed **agentic LLM workflows** for multi-step healthcare queries with tool discipline, structured reasoning, and grounding rules; response stability ~**25%**, hallucinations cut >**30%**.
- Shipped **predictive ML pipelines** (scikit-learn, XGBoost) for patient no-show and care engagement scoring; recall +**15-20%** on high-risk categories, precision held >**90%** via class weighting, stratified sampling, threshold calibration.
- Packaged ML/LLM inference as **FastAPI / Flask** services on **Docker** with structured logging and load simulation; post-deploy defects ~**30%**.
- EHR preprocessing (Pandas, NumPy); dataset reliability >**98%**, downstream instability -**40%**.
- HIPAA-conscious governance: de-identification, data lineage, evaluation audit trails, stakeholder-facing system-limitation docs.

Energy Solutions International (Emerson)

Jun 2022 - Apr 2023

Database & DevOps Performance Engineer (Intern)

Hyderabad, India - Oil & Gas ERP

- Optimized **T-SQL stored procedures** on a compliance-sensitive ERP; query time -**20%**, retrieval latency -**25%**.
- Built **Jenkins CI/CD pipelines** for schema updates and stored-procedure deployments; deploy errors -**30%**, release cycle -**35 to -40%**.
- Designed performance dashboards with **SQL Server DMVs + Grafana**; incident recurrence -**25%**, root-cause resolution -**30%**.
- Tuned index maintenance and partition-aware queries; reconciliation runtime -**15%**, deadlocks -**15%**.

Suvidha Foundation

Jan 2022 - Mar 2022

Machine Learning Engineer

Hyderabad, India - Nonprofit, Video Analytics, Applied NLP

- Built **transformer-based video summarization** over 5,000+ recorded sessions; hierarchical summarization and timestamp-aligned clip extraction cut manual review **60-70%**.
- AI-selected highlights aligned with human-curated moments ~**85%**.
- Implemented **document Q&A with RAG**; hallucinations ~**30%**, contextual accuracy >**85%**.
- Deployed as **Flask** API with a lightweight web interface for non-technical staff.

PROJECTS

Agentic Healthcare Claims Processing & Fraud Risk Intelligence System

5-stage multi-agent pipeline (Intake, Validation, Consistency, Duplicate, Risk Scoring) with schema-validated JSON contracts between agents to prevent cascading hallucinations. RAG-grounded CPT/ICD validation against vector-indexed policies. Audit-ready explainable risk scoring.
Python, LangChain, FastAPI, pgvector, GPT-4, schema contracts

Manga Lens - Multi-Provider AI Vision Chrome Extension

Real-time manga translation extension shipped to the Chrome Web Store. Multi-provider abstraction (4 AI vision providers) with per-provider payload handling for worst-case failure isolation. 29 sites supported across Manifest V3 service workers.
TypeScript, Manifest V3, Service Workers, OpenAI / Anthropic / Google / OpenRouter

Dream Decoder - Multimodal Creative Platform

End-to-end multimodal pipeline: text description, intermediate prompt transformation, grounded image generation. ~30% contextual alignment gain, ~25-30% first-pass image success.
FastAPI, React + Vite, Tailwind, Perplexity Sonar, Replicate

YOLOv8 Driver Drowsiness Detection

Real-time computer vision for driver drowsiness; sub-50ms inference on edge hardware. Walk-forward validation and class-weighted training for imbalanced safety data.
PyTorch, YOLOv8, OpenCV, ONNX

EDUCATION

- Master of Science, Computer Science** - Kent State University

Focus: Applied Machine Learning, NLP, Database Systems.
- Aug 2023 - May 2025
- Bachelor of Technology, Computer Science** - KL University

Founder / Lead of E-Farming platform (university project, pilot with 3 cooperatives).
- Jun 2019 - May 2023

CERTIFICATIONS

- Deep Learning Specialization - DeepLearning.AI (Coursera)
- 2024

SKILLS

- Languages:** Python, TypeScript, JavaScript, SQL (T-SQL, PostgreSQL, SQLite), Java, C++
- AI / ML:** PyTorch, scikit-learn, XGBoost, LangChain, LlamaIndex, OpenAI, Anthropic, Hugging Face Transformers, Diffusers, YOLOv8
- LLM / GenAI:** RAG, agentic workflows, structured outputs, prompt engineering, evaluation pipelines, guardrails, grounding, embeddings, vector search
- Backend:** FastAPI, Flask, REST, WebSockets, Docker, Celery, Playwright
- Frontend:** React, Next.js, Vite, Tailwind, Chrome Extensions (Manifest V3)
- Data:** PostgreSQL + pgvector, MongoDB, Redis, Pandas, NumPy, Snowflake
- DevOps:** Docker, Jenkins, GitHub Actions, Grafana, Prometheus, AWS (EC2, S3, Lambda)
- Domains:** Healthcare AI (HIPAA-conscious), Applied AI, Enterprise AI, Workflow Automation, Computer Vision, Multimodal