

Deepak Mallampati

APPLIED AI / GENERATIVE AI ENGINEER

(330) 554-0770 | deepakmallampati00@gmail.com | linkedin.com/in/deepak-mallampati |
US-based, F-1 OPT (open to sponsorship) | Open to relocation

PROFESSIONAL SUMMARY

AI/ML engineer and AI-native builder with 4+ years shipping production Generative AI and LLM systems. I move from idea to working prototype in days using AI coding tools (Claude Code), then harden prototypes into clean, scalable products - shipped MangaLens (live Chrome extension, 29+ sites) and career-ops (autonomous AI job-search pipeline). Deep across the GenAI building blocks the role centers on: prompt engineering, RAG, fine-tuning, agentic frameworks, evals, and guardrails, with full-stack delivery on FastAPI/React and cloud.

SKILLS

Generative AI & LLMs • RAG Architecture • Agentic Frameworks • Fine-tuning (LoRA/QLoRA) • Evals & Guardrails • AI Coding Tools (Claude Code) • Full-Stack (FastAPI/React) • Rapid Prototyping

Generative AI & LLMs: GPT-4o/GPT-5, Claude Sonnet, Gemini, LLaMA | LangChain/LangGraph | RAG | Prompt Engineering | Multi-Agent Systems | Fine-tuning (LoRA, QLoRA, PEFT) | Guardrails/Content Safety

AI Coding & Product: Claude Code, Codex, Gemini CLI | Rapid prototyping → production | Hugging Face Transformers

Evaluation & Retrieval: RAGAS, BERTScore | FAISS, Pinecone, Weaviate | Semantic & Hybrid Search | Cross-Encoder Re-ranking

Full-Stack & Cloud: Python, JavaScript/TypeScript | FastAPI, React | AWS, Azure, GCP | Docker, Kubernetes | CI/CD | MLflow, W&B

PROFESSIONAL EXPERIENCE

Progress Solutions

Jul 2025 - Present

AI Engineer

USA

- Architected production-grade **agentic LLM systems** on a conductor-subagent orchestration framework, decomposing complex objectives into context-scoped subagents - reducing token consumption by **42%** while sustaining task accuracy across multi-step workflows.
- Engineered **retrieval-augmented generation (RAG)** pipelines over large-scale healthcare corpora, combining dense vector retrieval with cross-encoder re-ranking to lift answer relevance and materially reduce hallucination.
- Established a structured **LLM evaluation and monitoring framework** (RAGAS, BERTScore, custom metrics) with latency/accuracy dashboards to benchmark iterations and surface regressions before release.
- Optimized inference cost and latency through prompt compression, semantic caching, and model routing, sustaining SLA targets while reducing per-query overhead.
- Built fault-tolerant automation infrastructure with scheduled batch orchestration and exponential-backoff retry logic, reducing failed production runs by **60%**.

Vanguard

Jan 2024 - Jan 2025

AI Engineer Intern

USA

- Developed **AI agents and backend services** for Vanguard's internal AI platform across **25+ AI agents** and workflows, integrating with production data pipelines and observability tooling.
- Optimized prompts and evaluated **GPT-4o, Claude Sonnet, and Gemini**, improving task success rates by **27%** and informing platform model-selection decisions.
- Built **12 APIs and microservices**, reducing p95 latency from 1.5s to **800ms** and supporting **100,000+ requests daily** across 20+ internal teams.
- Engineered safeguards and content controls for policy-compliant AI responses, reducing unsafe outputs by **42%** while maintaining quality.

Kent State University

Oct 2023 - Jan 2024

ML & Gen AI Research Assistant

USA

- Fine-tuned **Qwen3 via 4-bit QLoRA** on an NVIDIA H100 using a curated instruction dataset, conditioning on a six-stage narrative schema for controllable long-form generation.
- Built a **privacy-preserving ML pipeline** (k-anonymity, l-diversity, Laplace differential privacy), cutting re-identification risk from **3.38% to 0.32%** while retaining 99% of baseline predictive performance.
- Authored a peer-review-ready IEEE-format conference paper documenting methodology and results.

Emerson

Jun 2022 - Apr 2023

ML Trainee

India

- Built supervised regression/classification models for equipment-failure prediction, anomaly detection, and capacity forecasting on operational data.
- Fine-tuned **BERT and RoBERTa** for entity extraction and text classification, achieving **89% F1** on custom NER tasks.

PROJECTS

MangaLens [Live on Chrome Web Store](#)

Real-time AI translation of manga/webtoons across 29+ sites - a single inline action replacing a multi-step manual chore.

career-ops

Autonomous AI-driven job-search pipeline (built with Claude Code) that scans listings, scores fit, and generates tailored applications overnight without supervision.

Privacy-Preserving Clinical ML Pipeline

k-anonymity + l-diversity + differential privacy framework quantifying the privacy-utility tradeoff on sensitive patient data.

EDUCATION

M.S. in Computer Science - Kent State University, OH

2025

GPA: 3.70 / 4.0